



UNIVERSITY OF CAPE TOWN
IYUNIVESITHI YASEKAPA • UNIVERSITEIT VAN KAAPSTAD

Machine Learning

Assignment 2

Model Fitting, Cross-Validation, Regularisation and Image
Scaling.

Julian Albert
ALBJUL005



Department of Statistical Sciences
University of Cape Town
South Africa
September 30, 2019

Contents

Problem 1	2
i) Model Fitting	2
ii) Validation and Out-Of-Sample Errors	2
Problem 2	4
i) Plot Data and Simulated Model	4
ii) Regularisation	5
iii) Cross-Validation	6
Problem 3	8
i) Image Scaling	8
ii) Eigenfaces	9
iii) Image Reconstruction	11

Problem 1

In this Problem I consider the model

$$y_i = 0.8x_i + \epsilon_i; \quad -1 \leq x \leq 1, i = 1, \dots, N$$

where: $\epsilon_i \sim \text{Normal}(0, 1)$

i) Model Fitting

Suppose $x \sim U(-1, 1)$, we can simulate a dataset of size $N = 30$ and fit the models below

$$g_1(x) = 0.5 + b_1x \quad (1)$$

$$g_2(x) = -0.5 + b_2x \quad (2)$$

These models are fit using the `lm()` function in R, the resulting model fits are depicted in Figure (1a). We can see that Model 1 lies above the true model whilst Model 2 lies below, this is fairly obvious as the intercepts are equidistant above/below zero. From Figure (1b) we obtain the mean models over 1000 fits, we may expect that Model 1 overestimates whilst Model 2 underestimates the true model.

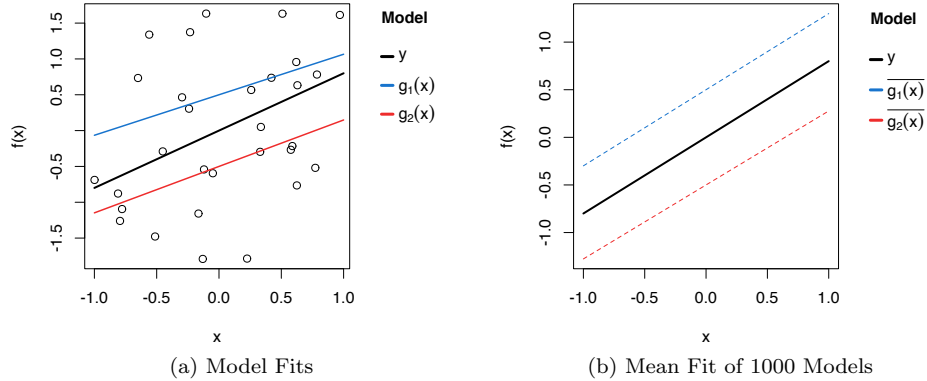


Figure 1: Models $g_1(x)$ and $g_2(x)$ fit once (left) and fit 1000 times then averaged (right).

ii) Validation and Out-Of-Sample Errors

Here we will consider the effect validation size has on different errors. First we define our dataset to be \mathcal{D} containing x and y which can be subdivided into a training set \mathcal{D}_{train} and a validation \mathcal{D}_{val} . To choose the best model $g^*(x)$ we fit both models $g_1(x_{train})$ and $g_2(x_{train})$ as defined in Equation (1) and (2), where x_{train} is the data from \mathcal{D}_{train} . We use the coefficients from the fitted model (m) denoted $\hat{\beta}_m$ to obtain predictions for our validation set using

$$\hat{y}^m = \hat{\beta}_m^T x_{val}$$

We can find the validation error E_{val}^m using

$$E_{val}^m = \frac{1}{N} \sum_{i=1}^N (y_{val_i}^m - \hat{y}_{val_i}^m)^2$$

This is done for both model $m = 1$ and $m = 2$, the model (m) chosen is the one which corresponds to the lower E_{val}^m . The chosen model is then $g^*(x)$ which is used to determine E_{out} such that it is the MSE between the true model $y = 0.8x$ and the predictions $\hat{y} = \beta^*x$ plus σ^2 .¹ We are interested in how E_{val} and E_{out} are impacted by the validation set size, to determine this we can first see how E_{val} fairs.

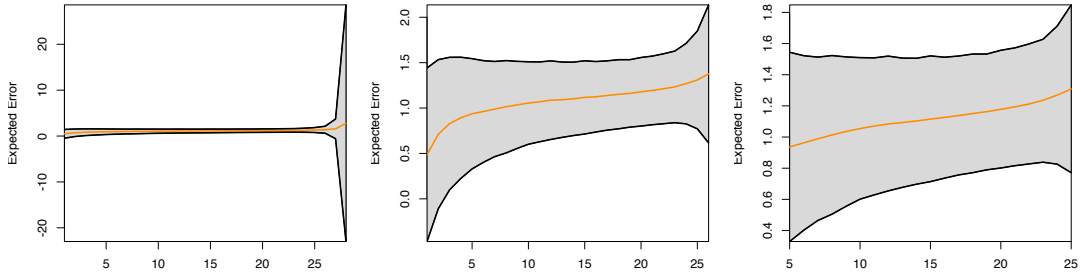


Figure 2: $E_{val} \pm \sigma_{val}$ for Different Validation Size Ranges.

From Figure (2) we see that if we use all the data as a validation set then the the uncertainty in E_{val} as measured by σ_{val} goes to extreme values (left-most plot). As the range is decreased (middle plot) we begin to see that when the validation set size is neither too small nor too large the uncertainty is more stable, thus E_{val} may provide a good estimate of E_{out} .

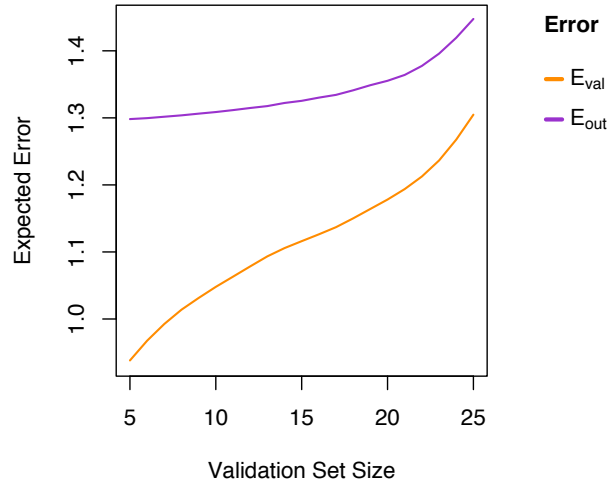


Figure 3: E_{val} and E_{out} for Validation Size = 5, 6, ..., 25 Averaged over 10000 Datasets.

¹ β^* represents the coefficients obtained from the training process of the model with the lowest E_{val}^m

Figure (3) illustrates the results of simulating 10 000 datasets of size $N = 30$, for each dataset \mathcal{D} we subdivide it into $\mathcal{D}_{train}^{(-i)}$ and $\mathcal{D}_{val}^{(i)}$ where $\mathcal{D}_{train}^{(-i)}$ denotes the training set, a set excluding i observations and $\mathcal{D}_{val}^{(i)}$ denotes the validation set, a set only containing i observations. To be explicit, together $\mathcal{D}_{train}^{(-i)}$ and $\mathcal{D}_{val}^{(i)}$ contain all observations in \mathcal{D} . We consider $i = 5, 6, \dots, 25$, the result is for each dataset we have training and validation sets of different sizes, this is done in the attempt to see how validation size impacts expected errors.

We can see that E_{val} tends to represent a conservative estimate of E_{out} , when we set aside more data for validation there are fewer training data points hence the expected validation error increases [1]. As the validation set size approaches the dataset size the two measure tend towards each other, this is due to there being no learning in the model.

Problem 2

Suppose $x \sim U(-1, 1)$, we can simulate a dataset of size $N = 30$ and fit the models below

$$y_i = \sin(\pi x_i) + \epsilon_i; \quad -1 \leq x \leq 1, i = 1, \dots, N \quad (3)$$

where: $\epsilon_i \sim \text{Normal}(0, 1)$

i) Plot Data and Simulated Model

Figure (4) shows the simulated data and true underlying model $y_i = \sin(\pi x_i)$.

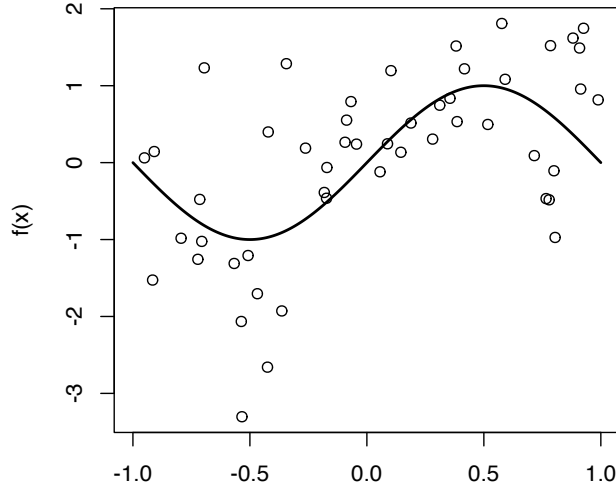


Figure 4: Simulated Model as per Equation (3).

ii) Regularisation

We can consider the model

$$y_i = \sum_{q=0}^{10} \beta_q L_q(x) \quad (4)$$

where $L_q(x)$ is a Legendre polynomial of order q . To fit this polynomial model is a special case of fitting a linear model in a new space \mathcal{Z} . This is done under a non-linear transformation $\Phi : \mathcal{X} \rightarrow \mathcal{Z}$, where for the Q^{th} order polynomial, Φ transforms vector \vec{x} into a \vec{z} of Legendre polynomials, thus

$$\vec{z} = [1 \quad L_1(x) \quad L_2(x) \quad \dots \quad L_{Q_f}(x)]$$

In matrix notation we can define $\vec{y} = \mathbf{Z}\vec{w} + \vec{\epsilon}$ where;

$$\mathbf{Z} = \begin{bmatrix} \vec{z}_1 \\ \vec{z}_2 \\ \vdots \\ \vec{z}_N \end{bmatrix} = \begin{bmatrix} 1 & L_1(x_1) & \dots & L_Q(x_1) \\ 1 & L_1(x_2) & \dots & L_Q(x_2) \\ \vdots & \vdots & & \vdots \\ 1 & L_1(x_N) & \dots & L_Q(x_N) \end{bmatrix}$$

The optimal solution for an unconstrained weight vector is given by

$$\hat{\vec{w}} = \left(\mathbf{Z}'\mathbf{Z} \right)^{-1} \mathbf{Z}'\vec{y} \quad (5)$$

Thus for a constrained problem (including a regularisation parameter λ) we get

$$\hat{\vec{w}} = \left(\mathbf{Z}'\mathbf{Z} + \lambda \mathbf{I} \right)^{-1} \mathbf{Z}'\vec{y} \quad (6)$$

Where \mathbf{I} is the $(Q_f + 1)$ identity matrix to penalise each coefficient by λ . This way complexity is controlled, we can see from Figure (5) that with $\lambda = 0$ there is a complex model fit, this is likely overfitting our data and will not generalise well to unseen data. However, for $\lambda = 5$ we see a curve that lies closer to the true model, although this is likely to perform worse in-sample the generalisation should perform better out-of-sample. The extreme case of λ being very large will simply fit a straight line through our data and is basically a naive fit/guess.

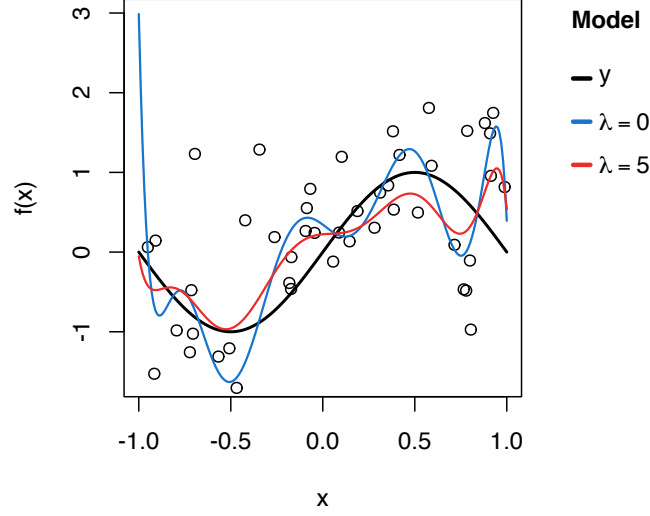


Figure 5: $L_{10}(x)$ Polynomial Fitted Models with Regularisation Parameter λ .

iii) Cross-Validation

So how do we choose the appropriate value for regularisation? We can use 10-fold cross-validation. This creates a quasi-OOS test of regularisation performance. We take our data and split it into training (IS) and testing (OOS) sets, we then get an OOS error based on the predictions between the true testing set response and our predicted response from the test set. The K-fold algorithm pseudo-code is shown in Algorithm (1).

Algorithm 1 10-Fold Cross Validation

```

1: procedure KFOLD.CV( $x, \lambda$ )  $\triangleright$  The CV error for data ( $x$ ) and complexity parameter ( $\lambda$ )
2:   Define  $n$  Observation,  $K$  folds and  $Q_f$  order.
3:   Shuffle the Data and Split data into  $n/K$  groups
4:    $\mathbf{I} \leftarrow$  matrix  $\triangleright$  Initialise  $(Q_f + 1) \times (Q_f + 1)$  Identity matrix
5:   for  $fold \leftarrow 1, K$  do
6:     Set an index equal to  $fold$ .
7:
8:     Subset training data to be group not fold of the shuffled data.
9:     Denoting the predictors as  $\mathbf{X}_{-fold}$  and response as  $\vec{y}_{-fold}$ 
10:
11:     Subset testing data to be group fold of the shuffled data.
12:     Denoting the predictors as  $\mathbf{X}_{fold}$  and response as  $y_{fold}$ 
13:
14:     Calculate the  $\hat{w}$  coefficients as per Equation 6 using  $\mathbf{X}_{-fold}$  and  $\vec{y}_{-fold}$ 
15:     Calculate prediction  $\hat{f}(\mathbf{X}_{fold}) = \mathbf{X}_{fold}\hat{w}$ 
16:     Cross-Validation Error is  $CV_{fold} = \frac{1}{K} \sum_{i=1}^K (y_{fold_i} - \hat{f}(\mathbf{X}_{fold_i}))$ 
17:   end for
18: end procedure

```

Thus for each λ we will get K number of CV errors, we can take the average of this for each value of λ to get an average error measure and find the λ that corresponds to the minimum of that. We generate a vector of regularisation parameters ($0.1 \leq \lambda \leq 10$) and obtain 10 CV errors per λ .

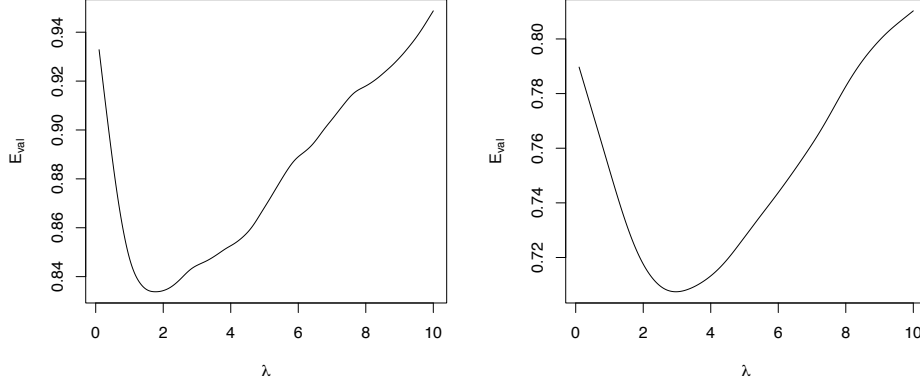


Figure 6: 10-fold Cross Validation Errors, Mean over folds (left) and median over folds (right).

From Figure (6) we can see that the over 10 folds, both the mean and median values show a fast initial decline in E_{val} before reaching some turning point — the point which acts as a proxy for where we begin to overfit. The median of CV per fold is plotted for sanity checking, but the minimum CV is taken to correspond with $\lambda^* \approx 1.7$. This λ^* is chosen and a model fit yielding Figure (7). We can see that the chosen model is more complex than $\lambda = 5$ in Figure (5), however, the fit seems reasonably close to the true model given the data.

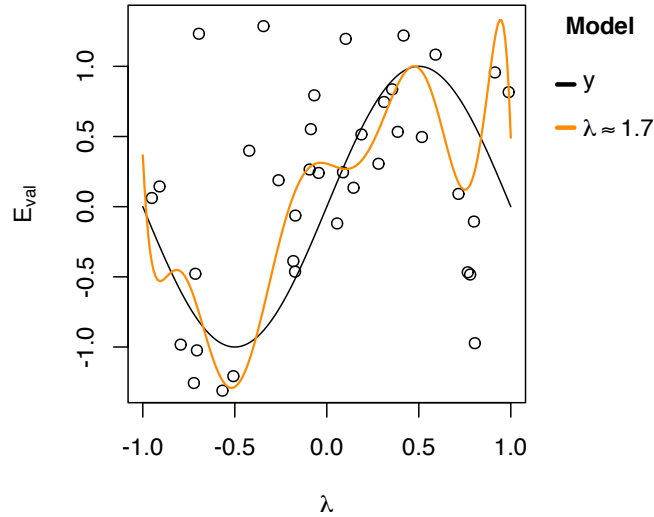


Figure 7: $L_{10}(x)$ Polynomial Fitted Models with Optimal Regularisation Parameter $\lambda \approx 2$ From 10-fold Cross Validation.

Problem 3

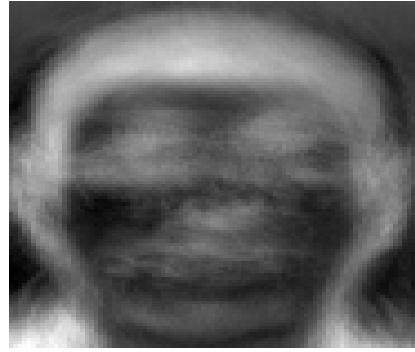
Here we will consider a dataset consisting of images from the `faces.zip` folder given to us. The folder contains $N = 400$ grey scale images which have dimension (width×height) 92×112 or $p = 10,304$.

i) Image Scaling

First we convert each image into a matrix of grey colour values using `image@grey` command in R. We correctly orient each of the images and collapse this matrix into a vector which is subsequently bound with all other image vectors to give a big matrix ($N \times Width * Ht$) representing all image colour values that has dimension (400×10304). As the rows represent each image we take the column means and column standard deviation to find the mean image (μ) and standard deviation (σ) image vector which is then “squashed” into a matrix.



(a) Mean Image



(b) Standard Deviation Image

Figure 8: Mean and Standard Deviation Images for 400 Images from `faces.zip`

We can use these two statistic matrices to scale our images, where the scaled image \mathbf{Z} is represented by

$$\mathbf{Z} = \frac{\mathbf{X} - \mu}{\sigma} \quad (7)$$

The resulting original vs scaled image `168.pgm` is plotted in Figure (9), we can see that the main difference is dark detail has been scaled to lighter shades of grey.



Figure 9: Original (\mathbf{X}) and Scaled (\mathbf{Z}) Image for File 168.pgm

ii) Eigenfaces

Eigenfaces are the set of eigenvectors used in human face recognition problems. The idea is to reduce the dimensions from (Height \times width) to a space where only the detail required to reproduce the image is kept. For basic facial recognition the dimensions need not be large as we can recognise a face as being a certain shape with eyes, a nose, a mouth etc. It is only when we need a higher level of detail that we would need more dimensions in our problem, reducing the dimension saves computational cost. In other words, as images of faces are very highly correlated we can compress the data to a low-dimensional subspace.

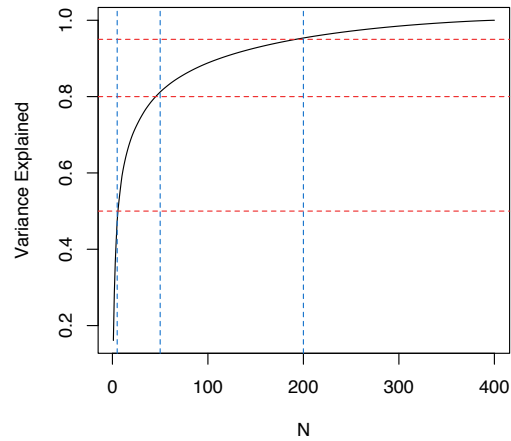


Figure 10: % Variance Explained as a Function of K Eigenfaces.

To determine how many eigenvectors are required we turn to Figure (10). Before reduction our covariance matrix (Σ) is $10,304 \times 10,304$ thus yielding 10,304 eigenvectors. In cases where

$N \ll p$ we can use an alternative definition for the covariance matrix. Let

$$\Sigma' = \frac{1}{N} \mathbf{X}' \mathbf{X}$$

represent a matrix \mathbf{A} which is $N \times N$. Now

$$\left(\frac{1}{N} \mathbf{X} \mathbf{X}' \right) \psi = \lambda \psi$$

Where ψ is $\frac{1}{N} \mathbf{X} \mathbf{X}'$'s eigenvectors. We now have

$$\begin{aligned} \Rightarrow \mathbf{X}^T \left(\frac{1}{N} \mathbf{X} \mathbf{X}' \right) \psi &= \lambda (\mathbf{X}^T \psi) \\ \Rightarrow \left(\frac{1}{N} \mathbf{X}^T \mathbf{X} \right) (\mathbf{X}^T \psi) &= \lambda (\mathbf{X}^T \psi) \\ \Rightarrow \phi = (\mathbf{X}^T \psi) &\text{ the eigenvector of } \Sigma_{p \times p} = \frac{1}{N} \mathbf{X}^T \mathbf{X} \end{aligned}$$

We know that the eigenvalues of the covariance matrix represent the variances of principal dimensions of our centered data, thus from Figure (10) shows the number of eigenfaces (5, 50, 200) explain about (50, 80, 95)% of the variance respectively. Choosing the number of eigenfaces is thus a function of how detailed we would like our image to be.



Figure 11: First 10 Eigenfaces.

Figure (11) shows the first 10 eigenfaces, we can see that the general shape of a face with eyes a nose and mouth are captured relatively early, however, 10 eigenfaces is not enough to identify people from the dataset. As the eigenfaces increase so does the image detail.

iii) Image Reconstruction

Here we can use the above derivations of eigenfaces along with the scaled image 115.pgm to plot the reconstructed versions of the image when the first 5, 50 and 200 eigenfaces are used. From Figure (10) we know that these respectively explain (50, 80, 95)% of the variance, further supported by Figure (12). We can clearly see that when the first 5 eigenfaces are used we only have a rough outline of a face, we are unclear of gender or hairstyle but can nearly identify the image is likely a human face. Using the first 50 eigenfaces shows a clearer image, eyes are open with longish hair, distinctive nose with a concave mouth shape — overall the level of detail is sufficient for identifying a human face. Lastly when we use the first 200 eigenfaces (explaining 95% of the variance) we get an image close to the target (scaled image). Distinct features such as teeth, dimples and ears are visible. Depending on the objective for the image use we can determine that this is enough to uniquely identify this person.



(a) $E_v = 5$



(b) $E_v = 50$



(c) $E_v = 200$



(d) Scaled Image

Figure 12: Scaled Image for File 115.pgm and Reconstructed Versions of the Image using Different Numbers of First Eigenfaces (E_v)

- END -

References

- [1] ABU-MOSTAFA, Y. S. *Learning from data*, vol. 4.

Appendix

R Code

```
1 ##
2 #
3 # Author: Julian Albert
4 # Date: 22 September 2019
5 #
6 # Description:
7 # ML Assignment 2 - Regularisation, validation set size and image eigenfaces.
8 #
9 #-----#
10
11 # 0. Clean Workspace, Directory and Library ----
12
13 ## Clean Workspace
14 rm(list=ls())
15 dev.off() # Close Plots
16 setwd("~/") # Clear Path to User
17
18 ## Locations
19 project_folder <- "/Documents/UCT/Coursework/MachineLearning"
20 loc_script <- "/Assignment_2/UCT_Assignment/Code/MachineLearning_Ass2"
21 loc_figs <- "/Assignment_2/UCT_Assignment/Figs"
22 loc_data <- "/Assignment_2/UCT_Assignment/Code/MachineLearning_Ass2/Faces"
23
24 ## Data File name
25 dat_faces_file <- numeric()
26 for(i in 1:400) dat_faces_file[i] <- paste(i, ".pgm", sep = "")
27
28 ## Directories
29 dir_script <- paste("~/", project_folder, loc_script, sep = '')
30 dir_figs <- paste("~/", project_folder, loc_figs, sep = '')
31 dir_data <- paste("~/", project_folder, loc_data, sep = '')
32
33 ## Set Working Directory to Script Location
34 setwd(dir_script)
35
36 ## Libraries - Lazyload
37 if (!require("pacman")) install.packages("pacman")
38 p_load(tidyverse, data.table, Cairo, parallel, doParallel, pixmap, rARPACK)
39
40 colour_vec <- c("dodgerblue3", "firebrick2", "forestgreen", "gold",
41                "darkorange", "darkorchid3") # colour for pretty plot
42
43 # Problem 1 ----
44
45 # Simulate a Dataset
46
47 set.seed(123)
48 N <- 30
49 x_seq <- seq(-1, 1, length.out = N)
50 epsilon <- rnorm(N, 0, 1)
51 x_dat <- runif(N, -1, 1)
52 y_dat <- 0.8*x_dat + epsilon
53 data <- as.data.frame(cbind(x_dat, y_dat))
54
55 # fit models g1 and g2
56 g1_betas <- lm(y_dat ~ 0 + x_dat, offset = rep(+0.5, length(x_dat)))$coefficients
57 g2_betas <- lm(y_dat ~ 0 + x_dat, offset = rep(-0.5, length(x_dat)))$coefficients
58 g1_pred <- +0.5 + g1_betas*x_seq
59 g2_pred <- -0.5 + g2_betas*x_seq
```

```

60
61 ## i) plot models and fits
62 setwd(dir_figs)
63 cairo_pdf("ML_Ass2_fig_P1_modelfits.pdf", height = 5, width = 5)
64 par(mar = par()$mar + c(0,0,0,5), pty = 's') # larger margins so legend on side
65 plot(data$y_dat~data$x_dat, xlab = "x", ylab = "f(x)")
66 lines(x_seq, 0.8*x_seq, col = 'black', lwd = 2)
67 lines(g1_pred ~ x_seq, col = colour_vec[1], lwd = 1.5)
68 lines(g2_pred ~ x_seq, col = colour_vec[2], lwd = 1.5)
69 legend("topright", title = expression(paste(bold("Model"))),
70       inset = c(-0.35, 0), lwd = 3, cex = 1, xpd = TRUE, bty = "n",
71       legend = c(expression(y), expression(g[1](x)), expression(g[2](x))),
72       col = c('black', colour_vec[1:2]), y.intersp = 1.5, x.intersp = 0.4,
73       seg.len = 0.8)
74 dev.off()
75
76 # from the plot data seems to have upper and lower band, both models same?
77
78 ## i) What if we fitted model 1000 times?
79 N_dat <- 30
80 N_sims <- 1000
81 x_seq <- seq(-1, 1, length.out = N_dat)
82 epsilon <- rnorm(N_dat, 0, 1)
83 x_dat <- runif(N_dat, -1, 1)
84 y_dat <- 0.8*x_dat + epsilon
85 data <- as.data.frame(cbind(x_dat, y_dat))
86 model <- 0.8*x_seq
87
88 g1_mat_pred <- matrix(NA, N_sims, N_dat)
89 g2_mat_pred <- matrix(NA, N_sims, N_dat)
90
91 for(i in 1:N_sims){
92
93   x_dat <- runif(N_dat, -1, 1)
94   epsilon <- rnorm(N_dat, 0, 1)
95   y_dat <- 0.8*x_dat + epsilon
96
97   # fit models g1 and g2
98   g1_betas <- lm(y_dat ~ 0 + x_dat, offset = rep(+0.5, length(x_dat)))$coefficients
99   g2_betas <- lm(y_dat ~ 0 + x_dat, offset = rep(-0.5, length(x_dat)))$coefficients
100   g1_mat_pred[i, ] <- +0.5 + g1_betas*x_seq
101   g2_mat_pred[i, ] <- -0.5 + g2_betas*x_seq
102
103 }
104
105 g1_pred_means <- colMeans(g1_mat_pred)
106 g2_pred_means <- colMeans(g2_mat_pred)
107
108 cairo_pdf("ML_Ass2_fig_P1_mean_modelfits.pdf", height = 5, width = 5)
109 par(mar = par()$mar + c(0,0,0,5), pty = 's') # larger margins so legend on side
110 plot(x_seq, model, lwd = 2, type = 'l', xlab = "x", ylab = "f(x)",
111      ylim = c(min(g2_pred_means), max(g1_pred_means)))
112 lines(x_seq, g1_pred_means, lty = "dashed", col = colour_vec[1])
113 lines(x_seq, g2_pred_means, lty = "dashed", col = colour_vec[2])
114 legend("topright", title = expression(paste(bold("Model"))),
115       inset = c(-0.35, 0), lwd = 3, cex = 1, xpd = TRUE, bty = "n",
116       legend = c(expression(y), bquote(bar(g[1](x))), bquote(bar(g[2](x)))),
117       col = c('black', colour_vec[1:2]), y.intersp = 1.5, x.intersp = 0.4,
118       seg.len = 0.8)
119 dev.off()
120
121 ## ii)
122
123 # Simulate 10'000 datasets of y size = 30
124 n_datasets <- 10000
125 x_seq <- seq(-1, 1, length.out = 1000)
126 true_model <- 0.8*x_seq
127
128 func.sim_dataset <- function(n_datasets, N)
129 {
130   # initialise data storage
131   data_df <- list()
132   # gen x, y

```

```

133   for(i in 1:n_datasets){
134     x_dat <- runif(N, -1, 1)
135     epsilon <- rnorm(N, 0, 1)
136     y_dat <- 0.8*x_dat + epsilon
137     data_df[[i]] <- data.frame(X = x_dat, Y = y_dat)
138   }
139   return(data_df)
140 }
141
142 dat_datasets <- func.sim_dataset(n_datasets, N)
143 ## Calc Eval and Eout
144 func.Eval_Eout <- function(data, i_start, i_end, x_seq, true_model)
145 {
146   ## Initialise Storage...
147   ## Need a matrix of Eval and Eout for each i
148   mat_of_errors <- matrix(0, nrow = length(i_start:i_end), ncol = 2)
149
150   # For Dval of size 5, 6, ..., 25 and Dtrain of size 25, 24, ..., 5
151   for(i in i_start:i_end)
152   {
153     # Subset data into training and validation
154     idx <- sample(1:NROW(data), i, replace = FALSE)
155     data_validation <- data[idx, ]
156     data_train <- data[-idx, ]
157     # fit models g1 and g2
158     g1_betas <- lm(Y ~ 0 + X, offset = rep(+0.5, length(X)), data_train)$coefficients
159     g2_betas <- lm(Y ~ 0 + X, offset = rep(-0.5, length(X)), data_train)$coefficients
160     g1_pred <- +0.5 + g1_betas*data_validation$X
161     g2_pred <- -0.5 + g2_betas*data_validation$X
162     # Calculate 1/n * (Y - g(Y)_val)^2
163     MSE_g1 <- mean((data_validation$Y - g1_pred)^2)
164     MSE_g2 <- mean((data_validation$Y - g2_pred)^2)
165
166     E_val <- min(MSE_g1, MSE_g2)
167
168     # select model with lower MSE
169     if(MSE_g1 < MSE_g2){
170       gstar_betas <- g1_betas
171       gstar_pred <- +0.5 + gstar_betas*(x_seq)
172     } else {
173       gstar_betas <- g2_betas
174       gstar_pred <- -0.5 + gstar_betas*(x_seq)
175     }
176
177     # Calculate OOS error and store
178     E_out <- mean( (gstar_pred - true_model)^2 ) + 1# ?
179     mat_of_errors[(i-i_start+1), ] <- c(E_val, E_out)
180   }
181
182   # Name columns to not get confused by which measure is which
183   dat_erros_val_and_oos <- as.data.frame(mat_of_errors)
184   names(dat_erros_val_and_oos) <- c("E_val", "E_out")
185
186   return(dat_erros_val_and_oos)
187 }
188
189 ## Parallel, 21 linear models over 10'000 datasets = 210'000 linear models
190 system.time(
191   dat_Errors <- mclapply(dat_datasets, # list
192     func.Eval_Eout, i_start = 5, i_end = 25,
193     x_seq = x_seq, true_model = true_model, # function & pars
194     mc.cores = detectCores()) # how many cores
195 )
196
197 # Select Eval and Eout to plot
198 Evals <- lapply(dat_Errors, function(x) x$E_val)
199 Eouts <- lapply(dat_Errors, function(x) x$E_out)
200 Evals_Mean <- bind_cols(Evals) %>% rowMeans()
201 Eouts_Mean <- bind_cols(Eouts) %>% rowMeans()
202
203 cairo_pdf("ML_Ass2_fig_P1_EvalEout.pdf", height = 5, width = 5)
204 par(mar = par()$mar + c(0,0,0,5), pty = 's') # larger margins so legend on side

```

```

206 plot(smooth.spline(y = Evals_Mean, x = 5:25), type = 'l', col = colour_vec[5],
207       ylim = c(min(Evals_Mean), max(Eouts_Mean)), lwd = 1.5,
208       xlab = "Validation Set Size", ylab = "Expected Error")
209 lines(smooth.spline(y = Eouts_Mean, x = 5:25), col = colour_vec[6], lwd = 1.5)
210 legend("topright", title = expression(paste(bold("Error"))),
211       inset = c(-0.35, 0), lwd = 3, cex = 1, xpd = TRUE, bty = "n",
212       legend = c(expression(E[val]), expression(E[out])),
213       col = colour_vec[5:6], y.intersp = 1.5, x.intersp = 0.4,
214       seg.len = 0.8)
215 dev.off()
216
217 # Why between 5:25? what happens over full range cant have val of size 30 in lm()
218 system.time(
219   dat_Errors <- mclapply(dat_datasets, # list
220     func.Eval_Eout, i_start = 1, i_end = 29,
221     x_seq = x_seq, true_model = true_model, # function & pars
222     mc.cores = detectCores()) # how many cores
223 )
224
225 ## What happens to uncertainty around Eval?
226 Evals <- lapply(dat_Errors, function(x) x$E_val)
227 Evals_Mean <- bind_cols(Evals) %>% rowMeans()
228 Evals_Mean <- Evals_Mean[-29]
229 Evals_sd <- bind_cols(Evals) %>% apply(1, function(x) sd(x))
230 Evals_sd <- Evals_sd[-29]
231 cairo_pdf("ML_Ass2_fig_P1_EvalwSigma_128.pdf", height = 5, width = 5)
232 plot(x = c(1, 28), y = c(min(Evals_Mean-Evals_sd), max(Evals_Mean+Evals_sd)), type = 'n',
233      xlab = "Validation Set Size", ylab = "Expected Error", yaxs = 'i', xaxs = 'i')
234 polygon(c(1:28, rev(1:28)), c( (Evals_Mean-Evals_sd), rev(Evals_Mean+Evals_sd) ),
235        col = "grey85")
236 lines(smooth.spline(y = Evals_Mean, x = 1:28), col = colour_vec[5], lwd = 2)
237 lines(smooth.spline(y = Evals_Mean+Evals_sd, x = 1:28), lwd = 2)
238 lines(smooth.spline(y = Evals_Mean-Evals_sd, x = 1:28), lwd = 2)
239 dev.off()
240
241 Evals <- lapply(dat_Errors, function(x) x$E_val)
242 Evals_Mean <- bind_cols(Evals) %>% rowMeans()
243 Evals_Mean <- Evals_Mean[-c(27, 28, 29)]
244 Evals_sd <- bind_cols(Evals) %>% apply(1, function(x) sd(x))
245 Evals_sd <- Evals_sd[-c(27, 28, 29)]
246 cairo_pdf("ML_Ass2_fig_P1_EvalwSigma_126.pdf", height = 5, width = 5)
247 plot(x = c(1, 26), y = c(min(Evals_Mean-Evals_sd), max(Evals_Mean+Evals_sd)), type = 'n',
248      xlab = "Validation Set Size", ylab = "Expected Error", yaxs = 'i', xaxs = 'i')
249 polygon(c(1:26, rev(1:26)), c( (Evals_Mean-Evals_sd), rev(Evals_Mean+Evals_sd) ),
250        col = "grey85")
251 lines(smooth.spline(y = Evals_Mean, x = 1:26), col = colour_vec[5], lwd = 2)
252 lines(smooth.spline(y = Evals_Mean+Evals_sd, x = 1:26), lwd = 2)
253 lines(smooth.spline(y = Evals_Mean-Evals_sd, x = 1:26), lwd = 2)
254 dev.off()
255
256 Evals <- lapply(dat_Errors, function(x) x$E_val)
257 Evals_Mean <- bind_cols(Evals) %>% rowMeans()
258 Evals_Mean <- Evals_Mean[-c(1:4, 26:29)]
259 Evals_sd <- bind_cols(Evals) %>% apply(1, function(x) sd(x))
260 Evals_sd <- Evals_sd[-c(1:4, 26:29)]
261
262 cairo_pdf("ML_Ass2_fig_P1_EvalwSigma_525.pdf", height = 5, width = 5)
263 plot(x = c(5, 25), y = c(min(Evals_Mean-Evals_sd), max(Evals_Mean+Evals_sd)), type = 'n',
264      xlab = "Validation Set Size", ylab = "Expected Error", yaxs = 'i', xaxs = 'i')
265 polygon(c(5:25, rev(5:25)), c( (Evals_Mean-Evals_sd), rev(Evals_Mean+Evals_sd) ),
266        col = "grey85")
267 lines(smooth.spline(y = Evals_Mean, x = 5:25), col = colour_vec[5], lwd = 2)
268 lines(smooth.spline(y = Evals_Mean+Evals_sd, x = 5:25), lwd = 2)
269 lines(smooth.spline(y = Evals_Mean-Evals_sd, x = 5:25), lwd = 2)
270 dev.off()
271
272 # Problem 2 ----
273
274 ## Simulate dataset
275 set.seed(123)
276 N = 50
277 x_seq <- seq(-1, 1, 0.01)
278 x_dat <- runif(N, -1, 1)

```



```

279 epsilon <- rnorm(N, 0, 1)
280 y_dat <- sin(pi*x_dat) + epsilon
281 data <- as.data.frame(cbind(x_dat, y_dat))
282 model <- sin(pi*x_seq)
283 Qf <- 10
284
285 setwd(dir_figs)
286 cairo_pdf("ML_Ass2_fig_P2_sim_model_init.pdf", height = 5, width = 5)
287 plot(x_seq, model, type = 'l', lwd = 2, xlab = "x", ylab = "f(x)",
288      ylim = c(min(y_dat), max(y_dat)))
289 points(data)
290 dev.off()
291
292 ## Legendre polynomial function
293 func.Legendre <- function(x, q)
294 {
295   Lq = 0
296   for(i in 0:q){
297     Lq = Lq + ((x^i)*choose(q, i)*choose((q+i-1)/2, q))
298   }
299   return((2^q)*Lq)
300 }
301
302 ## Target function with q-th order Legendre Polynomial
303 func.rand_Legfunc <- function(x, q)
304 {
305   fx = 0
306   beta_vec = runif(q+1, -1, 1)
307   for(i in 0:q){
308     fx = fx + (func.Legendre(x, i)*beta_vec[(i+1)])
309   }
310   return(fx)
311 }
312
313 ### Takes X -> Z maps to Z... z = [1 L1 L2 ... LQ]'
314 ### Z = [z1 z2 ... zN]' where z1 = [1 1 ... 1]', z2 = [L1(x1) L1(x2) ... L1(x3)]'
315
316 ## Fit model
317 func.fit_model <- function(X, Y, Qf, lambda = 0)
318 {
319   # initialise Z
320   Z <- matrix(0, nrow = NROW(X), ncol = (Qf+1))
321   # Lq for each column
322   for(i in 0:Qf){
323     Z[, (i+1)] <- func.Legendre(X, i)
324   }
325   # calc w* and hence predict
326   I <- diag(NCOL(Z))
327   Hat_mat <- Z %*% solve(crossprod(Z) + lambda*I) %*% t(Z)
328   Beta_Hat <- solve(crossprod(Z) + lambda*I) %*% t(Z) %*% Y
329   Y_hat <- Z %*% Beta_Hat
330
331   return(list(Hat_mat = Hat_mat, Beta_Hat = Beta_Hat, Y_hat = Y_hat))
332 }
333
334 test0 <- func.fit_model(x_dat, y_dat, Qf, lambda = 0)
335 test5 <- func.fit_model(x_dat, y_dat, Qf, lambda = 5)
336
337 func.predict <- function(x_seq, betas, Qf)
338 {
339   Z <- matrix(0, nrow = NROW(x_seq), ncol = (Qf+1))
340   for(i in 0:Qf){
341     Z[, (i+1)] <- func.Legendre(x_seq, i)
342   }
343   pred <- Z %*% betas
344   return(pred)
345 }
346
347 pred0 <- func.predict(x_seq, test0$Beta_Hat, Qf)
348 pred10 <- func.predict(x_seq, test5$Beta_Hat, Qf)
349
350 setwd(dir_figs)
351 cairo_pdf("ML_Ass2_fig_P2_modelfits.pdf", height = 5, width = 5)

```

```

352 par(mar = par()$mar + c(0,0,0,5), pty = 's') # larger margins so legend on side
353 plot(x_seq, model, type = 'l', lwd = 2, xlab = "x", ylab = "f(x)",
354      ylim = c(min(pred0, pred10), max(pred0, pred10)))
355 points(data)
356 lines(x_seq, pred0, col = colour_vec[1], lwd = 1.5)
357 lines(x_seq, pred10, col = colour_vec[2], lwd = 1.5)
358 legend("topright", title = expression(paste(bold("Model"))),
359        inset = c(-0.35, 0), lwd = 3, cex = 1, xpd = TRUE, bty = "n",
360        legend = c(expression(y), expression(lambda==0), expression(lambda==5)),
361        col = c('black', colour_vec[1:2]), y.intersp = 1.5, x.intersp = 0.4,
362        seg.len = 0.8)
363 dev.off()
364
365 ## Cross validation - K-fold
366 func.cv <- function(lambda, data, k.folds)
367 {
368   ## is data a subset of the fold number?
369   stopifnot(NROW(data) %% k.folds == 0)
370   ## data shuffle all the rows
371   data <- as_tibble(data)
372   data_shuffled <- sample_frac(data, 1L)
373   MSE <- numeric()
374   ## Split data into k folds
375   folds <- rep(1:k.folds, nrow(data_shuffled)/k.folds)
376   data_in_folds <- split(data_shuffled, folds)
377
378   for(fold in 1:k.folds)
379   {
380     ## Use fold as validation, all else as testing
381     dat_train <- bind_rows(data_in_folds[-fold])
382     dat_validate <- bind_rows(data_in_folds[fold])
383     # fit model with specified lambda
384     beta_hat <- func.fit_model(dat_train$x_dat, dat_train$y_dat, 10, lambda)$Beta_Hat
385     yhat <- func.predict(dat_validate$x_dat, beta_hat, 10)
386     MSE[fold] <- mean((yhat - dat_validate$y_dat)^2)
387   }
388   # MSE for the 10 folds and specified lambda value
389   return(MSE)
390 }
391
392 # For different values of lambda
393 lambda_vec <- seq(0.1, 10, length.out = 500)
394 testcv <- mclapply(lambda_vec, func.cv, data = data, k.folds = 10,
395                   mc.cores = detectCores())
396
397 cairo_pdf("ML_Ass2_fig_P2_CVerrors.pdf", height = 5, width = 10)
398 par(mfrow = c(1, 2))
399 lapply(testcv, mean) %>% # mean over folds...
400   unlist %>%
401   smooth.spline(x = lambda_vec) %>%
402   plot(type = 'l', xlab = expression(lambda),
403        ylab = expression(E[val]))
404 lapply(testcv, median) %>% # median over folds...
405   unlist %>%
406   smooth.spline(x = lambda_vec) %>%
407   plot(type = 'l', xlab = expression(lambda),
408        ylab = expression(E[val]))
409 dev.off()
410
411 # approximate lambda based on plots
412 lambdaOpt <- 1.70
413 testOpt <- func.fit_model(x_dat, y_dat, Qf = 10, lambda = lambdaOpt)
414 predOpt <- func.predict(x_seq, testOpt$Beta_Hat, Qf)
415
416 setwd(dir_figs)
417 cairo_pdf("ML_Ass2_fig_P2_opt_fit.pdf", height = 5, width = 5)
418 par(mar = par()$mar + c(0,0,0,5), pty = 's') # larger margins so legend on side
419 plot(x_seq, model, xlab = bquote(lambda), ylab = bquote(E[val]), type = 'l',
420      ylim = c(min(predOpt), max(predOpt)))
421 points(data)
422 lines(x_seq, predOpt, col = colour_vec[5], lwd = 1.5)
423 legend("topright", title = expression(paste(bold("Model"))),
424        inset = c(-0.35, 0), lwd = 3, cex = 1, xpd = TRUE, bty = "n",

```

```

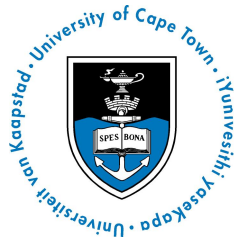
425     legend = c(expression(y), bquote(lambda%-%(lambdaOpt))),
426     col = c('black', colour_vec[5]), y.intersp = 1.5, x.intersp = 0.4,
427     seg.len = 0.8)
428 dev.off()
429
430 # Problem 3 ----
431
432 ## Initial pars, 400 92x112 images
433 N <- 400; Ht <- 112; Wdth <- 92
434 ## Read in images
435 images <- mclapply(1:N, function(i){
436   read.pnm(paste(dir_data, "/", dat_faces_file[i], sep = ""), cellres = 1)},
437   mc.cores = (detectCores() - 1))
438
439 # image(images[[1]]@grey, col = grey(seq(0, 1, length = 256)), axes = FALSE)
440
441 ## Rotate images to correct orientation
442 func.rotate <- function(x) t(apply(x, 2, rev))
443 image_oriented <- lapply(images, function(image){ func.rotate(image@grey) })
444 ## Oriented Images need to be in vector format
445 image_oriented_vec <- lapply(image_oriented, as.vector) %>% do.call("rbind", .)
446 ## Rows are image no. indices... column mean/sd gives pixel mean across all images
447 image_mean <- colMeans(image_oriented_vec) %>% matrix(., nrow = Wdth, ncol = Ht)
448 image_sd <- apply(image_oriented_vec, 2, sd) %>% matrix(., nrow = Wdth, ncol = Ht)
449
450 ## Plot mean and standard deviation image
451 setwd(dir_figs)
452 cairo_pdf("ML_Ass2_fig_P3_mean_image.pdf", height = 5, width = 5)
453 image(image_mean, col = grey(seq(0, 1, length = 256)), axes = FALSE)
454 dev.off()
455
456 cairo_pdf("ML_Ass2_fig_P3_sd_image.pdf", height = 5, width = 5)
457 image(image_sd, col = grey(seq(0, 1, length = 256)), axes = FALSE)
458 dev.off()
459
460 ## Scale images
461 image_scaled <- lapply(image_oriented,
462   function(image){ (image - image_mean) / image_sd })
463
464 ## Image 168... plot original (oriented) and scaled
465 n <- 168
466
467 cairo_pdf("ML_Ass2_fig_P3_orient_image168.pdf", height = 5, width = 5)
468 image(image_oriented[[n]], col = grey(seq(0, 1, length = 256)), axes = FALSE)
469 dev.off()
470
471 cairo_pdf("ML_Ass2_fig_P3_scaled_image168.pdf", height = 5, width = 5)
472 image(image_scaled[[n]], col = grey(seq(0, 1, length = 256)), axes = FALSE)
473 dev.off()
474
475 ## Want each row to be a width*ht vector
476 images_scaled_vec <- lapply(image_scaled,
477   function(image){ as.vector(t(image)) }) %>%
478   do.call("rbind", .)
479
480 ## Get Eigenvalues/Vectors using prcomp >> unit eigenvectors
481 ### N << p (much smaller) do eigen decomp on smaller matrix
482 ### rather use XXt
483 ## new eigen vector is Xt psi
484 A <- 1/N * images_scaled_vec %*% t(images_scaled_vec)
485 PCA <- prcomp(A)
486 image_eigen_vectors <- t(images_scaled_vec) %*% PCA$rotation
487 image_eigen_vectors_norm <- apply(image_eigen_vectors, 2, function(x){ x/sqrt(sum(x^2)) })
488
489 check_scree <- eigen(A)
490 check_eigs <- t(images_scaled_vec) %*% check_scree$vectors
491 setwd(dir_figs)
492 cairo_pdf("ML_Ass2_fig_P3_eigenvalues.pdf", height = 5, width = 5)
493 plot(cumsum(check_scree$values)/sum(check_scree$values), type = 'l',
494   xlab = "N", ylab = "Variance Explained")
495 abline(h = c(0.5, 0.8, 0.95), lty = 'dashed', col = colour_vec[2])
496 abline(v = c(5, 50, 200), lty = 'dashed', col = colour_vec[1])
497 dev.off()

```

```

498
499 ## Plot first ten eigenfaces
500
501 cairo_pdf("ML_Ass2_fig_P3_eigenfaces.pdf", height = 7.5, width = 10)
502 par(mfrow = c(2, 5))
503 par(mar = c(0.2, 0.2, 0.2, 0.2))
504 for (i in 1:10){
505     eigen_vec_mat <- matrix(image_eigen_vectors_norm[, i], nrow = 92, byrow = TRUE)
506     image(eigen_vec_mat, col = grey(seq(0, 1, length = 256)), axes = FALSE)
507 }
508 dev.off()
509
510 ## reconstruct an image based on eigenfaces used
511 func.reconstruct <- function(images_scaled_vec, image_eigen_vectors_norm)
512 {
513     ## Project onto the Eigen Space
514     project.eig <- t(data.matrix(images_scaled_vec)) %*% image_eigen_vectors_norm
515     ## Project back
516     reconstruct.img <- project.eig %*% t(image_eigen_vectors_norm)
517     ## Plot the image... Vector needs to be matrix
518     image(matrix(reconstruct.img, nrow = 92, byrow = TRUE),
519           col = grey(seq(0, 1, length = 256)), axes = FALSE)
520 }
521
522 ## Image 115, c(x, y) corresponds to number of eigenfaces used
523 n = 115
524 cairo_pdf("ML_Ass2_fig_P3_recon5_image115.pdf", height = 5, width = 5)
525 func.reconstruct(images_scaled_vec[n, ], image_eigen_vectors_norm[, c(1:5)])
526 dev.off()
527
528 cairo_pdf("ML_Ass2_fig_P3_recon50_image115.pdf", height = 5, width = 5)
529 func.reconstruct(images_scaled_vec[n, ], image_eigen_vectors_norm[, c(1:50)])
530 dev.off()
531
532 cairo_pdf("ML_Ass2_fig_P3_recon200_image115.pdf", height = 5, width = 5)
533 func.reconstruct(images_scaled_vec[n, ], image_eigen_vectors_norm[, c(1:200)])
534 dev.off()
535
536 cairo_pdf("ML_Ass2_fig_P3_scaled_image115.pdf", height = 5, width = 5)
537 image(image_scaled[[n]], col = grey(seq(0, 1, length = 256)), axes = FALSE)
538 dev.off()

```



Plagiarism Declaration Form

A copy of this form, completed and signed, to be attached to all coursework submissions to the Statistical Sciences Department.

COURSE CODE: **STA5068Z**
COURSE NAME: **Machine Learning**
STUDENT NAME: **Julian Albert**
STUDENT NUMBER: **ALBJUL005**
GROUP NUMBER: **1**

PLAGIARISM DECLARATION

- I know that plagiarism is wrong. Plagiarism is to use another's work and pretend that it is one's own.
- I have used a generally accepted citation and referencing style. Each contribution to, and quotation in, this tutorial/report/project from the work(s) of other people has been attributed, and has been cited and referenced.
- This tutorial/report/project is my own work.
- I have not allowed, and will not allow, anyone to copy my work with the intention of passing it off as his or her own work.
- I acknowledge that copying someone else's assignment or essay, or part of it, is wrong, and declare that this is my own work.
- Agreement to this statement does not exonerate me from the University's plagiarism rules.

Signature:

A handwritten signature in black ink, appearing to read 'Julian Albert', written over a horizontal line.

Date: September 30, 2019